

Understanding the Instruction Set Architecture: The Foundation of Computation

1. Introduction to Instruction Set Architecture (ISA)

The Instruction Set Architecture (ISA) serves as a fundamental concept in computer architecture, representing the critical interface between a computer's hardware and the software that runs on it. It is an abstract model that defines how the Central Processing Unit (CPU) is controlled by software, specifying the processor's capabilities and the mechanisms for invoking those capabilities.¹ The ISA is, in essence, a contract that dictates how a programmer, whether an assembly language programmer, a compiler writer, or an application developer, interacts with the hardware. It is the portion of the machine visible to these users, providing the sole means through which they can direct the hardware's operations.²

This interface abstracts away the intricate details of the underlying hardware implementation, such as the specific microarchitecture or circuit design. This abstraction is paramount because it allows for the evolution of hardware (e.g., faster, more power-efficient processors) without requiring changes to the software that conforms to that ISA. Conversely, different software can run on any hardware implementation that adheres to the same ISA. This decoupling fosters innovation in both hardware and software domains independently. For instance, any Arm-based processor, regardless of the specific manufacturer or core design, will execute software written to the Arm ISA specifications in the same way.²

The importance of the ISA extends to several aspects of computing. It defines:

- **Supported data types:** The kinds of data the processor can operate on (e.g., integers, floating-point numbers of various sizes).
- **Registers:** The internal storage locations within the CPU accessible by instructions.
- **Memory management:** How the hardware interacts with main memory, including features like virtual memory.
- **Instruction repertoire:** The set of operations the microprocessor can execute.
- **Input/Output model:** How the processor communicates with peripheral devices.²

An ISA can also be extended by adding new instructions, capabilities, or support for larger addresses and data values.² Understanding the ISA and how compilers utilize its instructions can empower developers to write more efficient code and to debug more

effectively by understanding compiler output.² The stability and definition provided by the ISA are crucial for the development and longevity of software ecosystems.

2. The Anatomy of Machine Instructions

Machine instructions are the most elementary commands that a computer's processor can execute. Each instruction is a sequence of bits that directs the CPU to perform a specific operation. The structure and interpretation of these instructions are defined by the ISA.

2.1. Core Components of a Machine Instruction

While the exact format varies significantly between different ISAs, machine instructions generally consist of several key fields³:

- **Opcode (Operation Code):** This is arguably the most crucial part of an instruction. The opcode is a pattern of bits that specifies the particular operation to be performed by the CPU, such as addition, subtraction, data movement (load, store, move), logical operations (AND, OR, NOT), or control flow changes (branch, jump, call).³ The length of the opcode field determines the maximum number of distinct operations an ISA can support; for example, an 8-bit opcode field allows for up to $2^8=256$ different operations.⁴
- **Operands:** These fields specify the data or the locations of data (addresses) that the operation will act upon. An instruction can have zero, one, two, or more operands.³ Operands can take several forms:
 - **Register Identifier:** Specifies one of the CPU's internal registers. The number of bits in this field limits the number of addressable registers.⁴
 - **Immediate Value:** The actual data value is embedded directly within the instruction itself.⁴
 - **Memory Address:** Specifies a location in main memory. This can be a direct address or an address that needs to be calculated using an addressing mode.⁴
- **Addressing Mode Specifiers:** Often, parts of the opcode or separate fields within the instruction indicate how the effective address of an operand is to be determined.³ This allows for flexibility in accessing data, as will be detailed in the next section.

The interplay between the opcode and operand fields defines the instruction's function. The processor decodes the opcode to understand what to do, and then uses the operand fields (interpreted according to any addressing mode information) to

fetch the necessary data and execute the operation.

2.2. Instruction Formats: Defining Structure by Operand Count

Instruction formats refer to the layout or structure of machine-level instructions, dictating how opcodes and operands are arranged. The number of explicit address fields in an instruction is a primary way to classify these formats, directly impacting program length and CPU complexity.³

- **Zero-Address Instructions:** These instructions do not explicitly specify any operands or addresses within the instruction itself. They are typically used in stack-based architectures. Operations are performed implicitly on the top one or two elements of a stack, and the result is pushed back onto the stack.³ For example, an ADD instruction would pop the top two values from the stack, add them, and push the sum.
 - *Example Expression:* $X = (A + B) * (C + D)$
 - *Zero-Address Sequence:* PUSH A; PUSH B; ADD; PUSH C; PUSH D; ADD; MUL; POP X.³ The conciseness of instructions in this format can lead to shorter programs in terms of instruction count for certain computations, but relies heavily on efficient stack management.
- **One-Address Instructions:** These instructions specify a single operand, which is often implicitly assumed to be an accumulator register within the CPU. The accumulator acts as a default source and destination for operations.³
 - *Format Example:* opcode | operand/address | mode.³
 - *Example Expression:* $X = (A + B) * (C + D)$
 - *One-Address Sequence (AC is accumulator):* LOAD A (AC=M[A]); ADD B (AC=AC+M); STORE T (M=AC); LOAD C (AC=M[C]); ADD D (AC=AC+M); MUL T (AC=AC*M); STORE X (M[X]=AC).³ This format simplifies instruction decoding but may require more instructions overall due to the need to constantly load and store values to/from the accumulator.
- **Two-Address Instructions:** These instructions specify two operands. One operand often serves as both a source and the destination for the result, or the operands can be two sources with the result overwriting one of them or being stored in a designated location if implied or specified by a separate field.³
 - *Format Example:* opcode | Destination Address | Source Address | mode.³
 - *Example Expression:* $X = (A + B) * (C + D)$
 - *Two-Address Sequence (R1, R2 are registers):* MOV R1, A (R1=M[A]); ADD R1, B (R1=R1+M); MOV R2, C (R2=M[C]); ADD R2, D (R2=R2+M); MUL R1, R2 (R1=R1*R2); MOV X, R1 (M[X]=R1).³ This format offers more flexibility than

one-address instructions and can reduce the number of memory accesses if registers are used effectively.

- **Three-Address Instructions:** These instructions specify three operands, typically two source operands and one destination operand.³ This format allows for more complex operations to be expressed in a single instruction and often maps more directly to expressions in high-level languages.
 - *Format Example:* opcode | Destination address | Source1 address | Source2 address | mode.³
 - *Example Expression:* $X = (A + B) * (C + D)$
 - *Three-Address Sequence (R1, R2 are registers, X is memory):* ADD R1, A, B ($R1 = M[A] + M$); ADD R2, C, D ($R2 = M[C] + M$); MUL X, R1, R2 ($M[X] = R1 * R2$).³ While potentially leading to shorter instruction sequences for complex calculations, three-address instructions are typically longer, requiring more bits to encode the three addresses. This format is common in many RISC architectures where operands are primarily registers.

The choice of instruction format is a fundamental ISA design decision. It reflects a trade-off between instruction length (and thus code density and memory bandwidth requirements), the complexity of the CPU's decoding and control logic, and the number of registers available. For instance, architectures with many registers (common in RISC) can effectively utilize two- or three-address formats where operands are registers, minimizing memory traffic. Conversely, architectures with fewer registers might rely more on one-address formats or two-address formats where one operand is a memory location.

2.3. Fixed vs. Variable Length Instructions

Another critical aspect of instruction design is whether instructions have a fixed length or a variable length.

- **Fixed-Length Instructions:** In this scheme, every instruction occupies the same number of bits (e.g., 32 bits in many RISC architectures like MIPS⁵).
 - **Advantages:**
 - **Simplified Fetch and Decode:** The CPU knows exactly how many bytes to fetch for the next instruction, simplifying the instruction fetch mechanism. Decoding is also generally faster and simpler because the locations of opcode and operand fields are consistent.⁵
 - **Easier Pipelining:** Predictable instruction boundaries facilitate efficient instruction pipelining, a key technique for improving processor

performance.⁵

- **Disadvantages:**
 - **Potential Wasted Space:** Simpler instructions that do not need all the available bits for operands or addresses might still occupy the full instruction word, leading to lower code density.⁵ For example, an instruction that only operates on registers might not need bits allocated for a large memory address, but those bits are still part of the fixed length.
 - **Limited Operand/Address Size:** The fixed length imposes constraints on the size of immediate values or the number of bits available for direct memory addresses within a single instruction.
- **Variable-Length Instructions:** In this scheme, instructions can have different lengths depending on the operation and the operands involved (e.g., CISC architectures like x86).
 - **Advantages:**
 - **Improved Code Density:** Instructions can be tailored to use only the number of bits necessary for their specific operation and operands. This can result in smaller program sizes, reducing memory usage and potentially improving cache performance.⁷
 - **Flexibility:** Allows for a wider range of operations and addressing modes to be encoded, including complex instructions with multiple memory operands.
 - **Disadvantages:**
 - **Complex Fetch and Decode:** The CPU must first determine the length of the current instruction before it can fetch and decode it completely, and then determine where the next instruction begins. This adds complexity to the control unit and can slow down instruction processing.⁸
 - **Difficult Pipelining:** The variable nature makes it harder to implement efficient instruction pipelines, as instruction boundaries are not fixed.

The choice between fixed and variable length instructions is a core differentiator between RISC and CISC philosophies. RISC architectures typically favor fixed-length instructions for their simplicity and performance benefits in pipelined execution⁵, while CISC architectures often use variable-length instructions to achieve higher code density and provide a richer, more powerful instruction set at the assembly level.⁷ However, modern processors often incorporate techniques to mitigate the disadvantages of their chosen approach; for example, some RISC ISAs (like ARM's

Thumb-2⁹⁾ introduce shorter instruction formats for common operations to improve code density, creating a hybrid approach.

3. Addressing Modes: Specifying Operand Locations

Addressing modes are the techniques employed by a CPU to determine the location of an operand referenced in an instruction.¹⁰ They provide rules for interpreting or modifying the address field of an instruction before the operand is actually accessed.¹⁰ The choice and variety of addressing modes significantly influence the flexibility and efficiency of an ISA, affecting how easily data structures like arrays, pointers, and records can be manipulated in assembly language.¹⁰

3.1. Purpose and Importance of Addressing Modes

The primary purpose of addressing modes is to provide versatile ways to specify operand locations, whether they are in registers, embedded within the instruction, or reside in memory. This flexibility is crucial for several reasons:

- **Efficient Data Handling:** They allow programmers and compilers to generate code that efficiently accesses various data structures. For example, indexed addressing is ideal for array element access, while register indirect addressing is useful for pointer dereferencing.¹⁰
- **Program Control:** Modes like PC-relative addressing are essential for implementing branches, jumps, and loops, enabling program control flow.¹⁰
- **Code Relocatability:** Certain addressing modes, such as PC-relative and base-register addressing, facilitate the creation of position-independent code, which can be loaded and executed at any memory address without modification.¹⁰ This is vital for modern operating systems and shared libraries.
- **Reduced Instruction Size:** By providing implicit ways to determine addresses or by using shorter address specifiers (like register numbers instead of full memory addresses), addressing modes can help reduce the overall size of instruction fields and, consequently, programs.¹⁰
- **Programming Convenience:** They can simplify assembly language programming by offering higher-level ways to think about data access, reducing the complexity of manual address calculations.¹⁰

3.2. Common Types of Addressing Modes

Numerous addressing modes exist, with specific ISAs implementing a subset tailored to their design philosophy. Some of the most common types include:

- **Implied / Implicit Addressing Mode:**
 - **Explanation:** The operand's location is inherent in the instruction's definition; no explicit address field is needed. This is common for instructions operating on specific registers (like an accumulator or stack pointer) or for zero-address instructions.¹⁰
 - **Example:** A CMA (Complement Accumulator) instruction implicitly operates on the accumulator.¹¹ PUSH and POP operations on a stack-based machine are also examples.
- **Immediate Addressing Mode:**
 - **Explanation:** The operand itself is part of the instruction. The value is directly embedded in the instruction's operand field.¹⁰
 - **Symbol:** Often denoted by #.¹⁰
 - **Example:** MOV AL, 35H (moves the hexadecimal value 35 into register AL).¹⁰
ADD R4, #3 (adds the constant 3 to register R4).⁵
 - **Purpose:** Useful for initializing registers or performing operations with constants. The size of the immediate value is limited by the number of bits allocated in the instruction format.
- **Register Direct Addressing Mode:**
 - **Explanation:** The operand is contained in a CPU register. The instruction's address field specifies which register holds the data.¹⁰
 - **Example:** MOV AX, CX (moves the content of register CX to register AX).¹⁰
ADD R4, R3 ($R4 \leftarrow R4 + R3$).⁵
 - **Purpose:** Provides very fast access to operands since registers are internal to the CPU. No memory access is required to fetch the operand.
- **Register Indirect Addressing Mode:**
 - **Explanation:** The instruction's address field specifies a CPU register, but this register does not contain the operand itself. Instead, it holds the memory address of the operand.¹⁰
 - **Example:** MOV AX, (BX) (moves the content of the memory location whose address is in register BX to register AX).¹⁰ ADD R4, (R1) ($R4 \leftarrow R4 + M$).⁵
 - **Purpose:** Essential for implementing pointers and accessing data structures where addresses are computed at runtime. Requires one memory reference to fetch the operand.
- **Direct Addressing / Absolute Addressing Mode:**
 - **Explanation:** The instruction's address field contains the actual memory address (effective address) of the operand.¹⁰
 - **Symbol:** Often denoted by [].¹⁰

- **Example:** ADD AL, [0301H] (adds the content of memory location 0301H to register AL).¹⁰ ADD R1, (1001) ($R1 \leftarrow R1 + M$).⁵
- **Purpose:** Simple way to access static data variables whose addresses are known at compile time. Requires one memory reference.
- **Indirect Addressing Mode (Memory Indirect):**
 - **Explanation:** The address field of the instruction specifies a memory location. This memory location, in turn, contains the effective address of the operand.¹⁰
 - **Symbol:** Often denoted by @ or ().¹⁰
 - **Example:** MOV AX, [[5000H]] (the memory location at 5000H contains the address of the actual data to be moved to AX).¹⁰
 - **Purpose:** Allows for complex pointer manipulations, such as pointers to pointers. Requires two memory references to fetch the operand, making it slower.
- **Indexed Addressing Mode:**
 - **Explanation:** The effective address of the operand is calculated by adding the content of an index register (which often holds an offset) to a base address (which can be part of the instruction or in another register).¹⁰
 - **Example:** MOV AX, (effective address = content of SI register + 5).¹⁰ ADD R4, 100(R1) ($R4 \leftarrow R4 + M$, where R1 is the base and 100 is a displacement/offset, a common form in MIPS also called Displacement addressing).⁵
 - **Purpose:** Extremely useful for accessing elements in arrays or tables, where the base address points to the start of the array and the index register holds the element's offset.
- **Base Register Addressing Mode:**
 - **Explanation:** Similar to indexed addressing, the effective address is formed by adding the content of a base register to a displacement value specified in the instruction.¹⁰ The base register typically points to the start of a data segment or record.
 - **Formula:** $EA = \text{Base register} + \text{Displacement}$.¹⁰
 - **Purpose:** Supports program relocation and access to fields within records or structures.
- **PC-Relative Addressing Mode:**
 - **Explanation:** The effective address is calculated by adding a displacement (offset) value, specified in the instruction, to the current value of the Program Counter (PC).¹⁰
 - **Formula:** $EA = PC + \text{Displacement}$.¹⁰
 - **Purpose:** Primarily used for implementing conditional and unconditional

branches. Since the offset is relative to the PC, the code can be position-independent.

- **Auto-Increment / Auto-Decrement Addressing Mode:**

- **Explanation:** The register used for indirect addressing is automatically incremented (after operand access) or decremented (before operand access) by the size of the operand.¹⁰
- **Example (Increment):** Add R1, (R2)+ ($R1 = R1 + M$; $R2 = R2 + d$).¹⁰
- **Example (Decrement):** Add R1, -(R2) ($R2 = R2 - d$; $R1 = R1 + M$).¹⁰
- **Purpose:** Very useful for stepping through arrays or implementing stacks (push/pop operations) efficiently in a loop.

- **Stack Addressing Mode:**

- **Explanation:** The operand is implicitly assumed to be at the top of the stack (or accessed relative to a stack pointer register).¹¹ Operations like PUSH and POP manipulate the stack and the stack pointer.
- **Purpose:** Fundamental for managing function calls, local variables, and expression evaluation in stack-based architectures or for stack operations in register-based machines.

The selection and implementation of addressing modes are critical design choices in an ISA. A rich set of addressing modes can simplify programming and improve code efficiency for certain tasks but can also increase the complexity of the CPU's control unit. RISC architectures, for instance, tend to offer fewer and simpler addressing modes compared to CISC architectures, relying more on register operations and explicit load/store instructions to access memory.⁵

4. RISC vs. CISC: Two Philosophies of ISA Design

The design of an Instruction Set Architecture is guided by fundamental philosophies that dictate the complexity and nature of its instructions. Two dominant philosophies have shaped processor design: Reduced Instruction Set Computing (RISC) and Complex Instruction Set Computing (CISC).⁷ These represent distinct approaches to optimizing processor performance, instruction encoding, and the hardware-software interface.

4.1. Core Philosophies and Design Goals

- **RISC (Reduced Instruction Set Computing):**

The core philosophy of RISC is to achieve higher performance by using a smaller, simpler set of instructions, where each instruction is designed to execute very

quickly, typically within a single clock cycle.⁷ The emphasis is on simplifying the hardware to make it faster. Complex operations are not performed by single complex instructions but are instead built up by sequences of these simple instructions, usually managed by a compiler. Key design goals include enabling efficient instruction pipelining and reducing the complexity of the control unit.⁷

- **CISC (Complex Instruction Set Computing):**

The CISC philosophy aims to improve performance and ease programming (especially at the assembly level) by providing a large, rich set of instructions, where individual instructions can perform complex, multi-step operations (e.g., a single instruction might load data from memory, perform an arithmetic operation, and store the result back to memory).⁷ The goal is to reduce the number of instructions required to complete a task, thereby potentially reducing the amount of memory needed for program code and the number of memory fetches for instructions.⁷

4.2. Key Characteristics and Architectural Differences

The differing philosophies lead to distinct architectural characteristics:

| Feature | RISC Architecture | CISC Architecture |
|-------------------------------|--|--|
| Instruction Set Size | Small, limited set of basic instructions. ⁸ | Large, extensive set of instructions. ⁸ |
| Instruction Complexity | Simple, single-cycle execution (ideally). ⁸ | Complex, multi-cycle execution for many instructions. ⁷ |
| Instruction Length | Typically fixed-length (e.g., 32-bit). ⁸ | Often variable-length. ⁸ |
| Addressing Modes | Fewer, simpler addressing modes. ⁸ | Many, complex addressing modes. ⁸ |
| Memory Access | Load/Store architecture: only | Instructions can often operate |

| | | |
|-----------------------------|---|---|
| | explicit load/store instructions access memory; operations are register-to-register. ⁶ | directly on memory operands. ⁵ |
| Registers | Large number of general-purpose registers. ⁸ | Fewer general-purpose registers, or specialized registers. ⁸ |
| Control Unit | Typically hard-wired control unit. ⁸ | Typically micro-programmed control unit. ⁸ |
| Instruction Decoding | Simple and fast. ⁸ | Complex and potentially slower. ⁸ |
| Pipelining | Easier to implement efficiently. ⁸ | More difficult to pipeline effectively due to instruction variability. ⁸ |
| Compiler Role | Heavy reliance on optimizing compilers to schedule instructions and manage resources. ⁵ | Compiler has less burden to break down high-level operations. |
| Code Density | Generally lower code density (more instructions per task). ⁷ | Generally higher code density (fewer instructions per task). ⁷ |
| Power Consumption | Often lower power consumption due to simpler hardware and fewer cycles per instruction on average. ⁷ | Can be higher due to complex decoding and multi-cycle instructions. ⁷ |
| Chip Size | Often smaller chip size for the core logic. ⁷ | Can be larger due to complex control logic. ⁷ |

Table based on information from.⁵

The hard-wired control units in RISC processors are faster because their logic is fixed

hardware, directly translating opcodes into control signals. In contrast, CISC processors often use micro-programmed control, where each machine instruction is interpreted by a sequence of micro-instructions (a microprogram) stored in a control ROM.⁸ This microcode adds a layer of interpretation, making it easier to implement complex instructions and fix bugs in the instruction logic, but it also introduces overhead, making individual instruction execution potentially slower.

The emphasis on a load/store architecture in RISC simplifies instruction design because arithmetic and logical operations do not need to handle memory operands directly. This means instructions do not need to specify memory addresses for data manipulation, only register numbers, leading to more compact and uniform instruction formats. CISC, by allowing operations directly on memory, can reduce the number of explicit load/store instructions, which was historically an advantage when memory was slow and compilers were less sophisticated.

4.3. Advantages and Disadvantages

RISC:

- **Advantages:**

- **Speed and Efficiency:** Simple instructions can be executed faster, often in a single clock cycle, and are well-suited for pipelining, leading to high throughput.⁷
- **Lower Power Consumption:** Simpler designs generally consume less power, making RISC ideal for mobile and embedded devices.⁷
- **Simpler Hardware Design:** Leads to smaller chip area, potentially lower manufacturing costs, and shorter design times.⁷

- **Disadvantages:**

- **Lower Code Density:** More simple instructions are typically needed to perform a complex task, leading to larger program sizes and potentially more instruction cache misses.⁷
- **Increased Memory Usage:** Larger number of instructions can consume more memory.⁷
- **Compiler Dependency:** Performance heavily relies on the quality of the compiler to optimize instruction scheduling and register usage.⁵
- **Limited Capabilities for Complex Tasks (in single instructions):** Some complex operations might be inefficient if they require very long sequences of simple instructions.⁷

CISC:

- **Advantages:**

- **Higher Code Density:** Complex instructions can perform more work, leading to smaller compiled code sizes, which can be beneficial for memory-constrained systems or reduce instruction fetch bandwidth.⁷
- **Simpler Compiler Design (Historically):** Compilers had an easier task mapping high-level language constructs to fewer, more powerful machine instructions.⁸
- **Memory Efficiency:** Fewer instructions for a given task can mean less memory traffic for fetching instructions.⁷

- **Disadvantages:**

- **Slower Execution (per instruction):** Complex instructions often take multiple clock cycles to execute, and the variability in execution time complicates pipelining.⁷
- **More Complex Hardware:** The control unit and decoding logic are more intricate, potentially leading to larger chip sizes and higher power consumption.⁷
- **Inefficient Use of Complex Instructions:** Studies have shown that compilers often tend to use simpler instructions from a CISC set, meaning the hardware complexity for rarely used complex instructions might not be justified.

4.4. Contextual Relevance and Modern Trends

The debate between RISC and CISC is not as clear-cut as it once was. Modern processor design often incorporates ideas from both philosophies. For example:

- Many CISC processors (like Intel's x86 series) internally break down complex CISC instructions into simpler, RISC-like micro-operations (μ ops) that are then executed by a high-performance, pipelined core. This allows them to retain backward compatibility with existing CISC software while leveraging RISC principles for internal execution.
- Some RISC architectures (like ARM) have added more complex instructions and features over time (e.g., SIMD extensions, Thumb-2 for variable-length instructions to improve code density⁹) to enhance performance for specific workloads and address traditional RISC weaknesses.

RISC architectures have become dominant in the mobile and embedded markets due to their power efficiency and performance per watt (e.g., ARM processors).⁷ They are also prevalent in high-performance computing in some segments. CISC architectures,

primarily x86, continue to dominate the desktop, laptop, and server markets due to a massive existing software ecosystem and continued performance improvements through advanced microarchitectural techniques.⁸

The choice between RISC and CISC, or a hybrid approach, depends on the specific design goals: target application domain, performance requirements, power constraints, cost, and the importance of backward compatibility. For embedded systems, RISC is often preferred due to its lower power consumption and smaller chip size.⁷ For general-purpose computing where legacy software is critical, CISC (or CISC with RISC-like cores) remains strong.

5. Comparative Analysis of ISAs: LC-3, MIPS, and ARM

To illustrate the concepts discussed, this section provides a comparative overview of three distinct ISAs: LC-3 (Little Computer 3), MIPS (Microprocessor without Interlocked Pipelined Stages), and ARM (Advanced RISC Machines). These ISAs represent different design points: educational simplicity, classic RISC, and commercially successful, evolving RISC, respectively.

5.1. LC-3 (Little Computer 3)

- Overview and Purpose:

The LC-3 is an ISA specifically designed for educational purposes, primarily to teach fundamentals of computer architecture and assembly language programming.¹³ Developed by Yale N. Patt and Sanjay J. Patel, it features a relatively simple instruction set yet is capable enough for moderately complex assembly programs and can even be a target for a C compiler.¹³ Its simplicity makes it an excellent tool for students to grasp core concepts without being overwhelmed by the complexities of commercial ISAs.¹² A simplified version of the LC-3 further reduces complexity for introductory learning, focusing on essential CPU operations and hardware design principles.¹²

- Key Architectural Features:

- **Word Size:** 16 bits for registers and data paths.¹²
- **Address Space:** 16-bit addressable memory, providing a 2¹⁶ location address space (word addressable).¹²
- **Registers:**
 - Eight 16-bit General-Purpose Registers (GPRs), R0-R7.¹²
 - A 16-bit Program Counter (PC).¹²
 - A 3-bit Processor Status Register (PSR) storing condition codes (N, Z, P -

Negative, Zero, Positive) based on the result of the last register write.¹²

- **Instruction Format:** All LC-3 instructions are 16 bits wide (fixed-length), a characteristic often associated with RISC principles.¹² The opcode is typically encoded in the 4 most significant bits (bits 15-12), with the remaining 12 bits for operands.¹²
- **Data Types:** Primarily operates on 16-bit two's complement integers. No native support for floating-point or unsigned arithmetic beyond what can be synthesized.¹³ I/O devices operate on ASCII characters.¹³
- **Control Unit:** Execution is regulated by a state machine, often implemented with a control ROM and microsequencing in more detailed models.¹³
- **Typical Instructions and Addressing Modes:**
 - **Instruction Set:** Implements around fifteen types of instructions in the standard LC-3.¹³ The simplified version further reduces this.
 - **Classification:** It is a load-store architecture; values in memory must be loaded into registers before being operated upon.¹³
 - **Arithmetic/Logical:** ADD, AND, NOT. These can use registers or sign-extended immediate values as operands.¹³ Subtraction can be implemented using NOT and ADD.
 - **Data Movement:** LD (Load Direct), LDR (Load Base+offset), LDI (Load Indirect), LEA (Load Effective Address), ST (Store Direct), STR (Store Base+offset), STI (Store Indirect).¹²
 - **Control Flow:**
 - BR (Conditional Branch): Branches based on the N, Z, P condition codes set by previous operations.¹²
 - JMP (Jump to address in register), RET (Return from subroutine, often a specific JMP).
 - JSR / JSRR (Jump to Subroutine, PC-relative or register-based, saves return address in R7).¹³
 - TRAP (System Call, for I/O and other OS services).¹³
 - **Addressing Modes (Simplified LC-3¹²):**
 - **PC-relative:** Address calculated by adding an offset to the (incremented) PC. Used for LD, ST, LEA, and BR.
 - **Register Indirect (Base+offset for LDR/STR):** Address calculated by adding an offset to the contents of a base register.
 - **Indirect:** Address read from a memory location whose address is PC-relative (LDI, STI).
 - **Register Direct:** Operands are in registers (for ADD, AND, NOT).

- **Immediate:** A small, sign-extended immediate value is part of the instruction for ADD and AND. The fixed 16-bit instruction length, load-store nature, and relatively simple operations classify LC-3 as RISC-like, especially in its intent and common implementations for teaching.¹²

5.2. MIPS (Microprocessor without Interlocked Pipelined Stages)

- Overview and Philosophy:
MIPS is a family of RISC ISAs originally developed by MIPS Computer Systems.⁶ It is a classic example of a RISC architecture, emphasizing a simple, streamlined instruction set designed for high performance through pipelining and compiler optimization.⁵ MIPS has been influential in academia and has seen use in embedded systems, networking equipment, and early game consoles.
- **Key Architectural Features (primarily MIPS I/MIPS32):**
 - **RISC Characteristics:** Load/store architecture, fixed-length instructions, a large number of GPRs, and simplified addressing modes.⁶
 - **Word Size:** Typically 32-bit (MIPS32) or 64-bit (MIPS64). MIPS I had 32-bit registers and addresses.⁶
 - **Registers (MIPS I/MIPS32):**
 - Thirty-two 32-bit GPRs (\$0 - \$31). Register \$0 is hardwired to zero. Register \$31 is often used as the link register for procedure calls.⁶
 - Special HI and LO registers for storing results of multiplication and division.⁶
 - A 32-bit Program Counter (PC).⁶
 - **Instruction Length:** All MIPS instructions are 32 bits long (fixed-length).⁵ This simplifies instruction fetch and decoding, aiding pipelining.⁵
 - **Instruction Formats:** Three primary formats⁶:
 - **R-type (Register):** opcode | rs | rt | rd | shamt | funct. Used for register-to-register operations.
 - **I-type (Immediate):** opcode | rs | rt | immediate. Used for operations with immediate values, loads, and stores.
 - **J-type (Jump):** opcode | address. Used for unconditional jumps.
 - **Data Types:** Supports 8-bit bytes, 16-bit half-words, and 32-bit words. Floating-point operations (if FPU coprocessor is present) support single (32-bit) and double-precision (64-bit) values.⁵
 - **Endianness:** Can be configured as either big-endian or little-endian (bi-endian).⁵

- **Typical Instructions and Addressing Modes:**

- **Instruction Categories:** Arithmetic, logical, data transfer (load/store), conditional branch, unconditional jump.
- **Arithmetic/Logical:** add, addu (add unsigned), sub, subu, addi (add immediate), and, or, xor, nor, andi, ori, xori, sll (shift left logical), srl, sra (shift right arithmetic).
- **Data Movement (Load/Store):** lw (load word), sw (store word), lb (load byte), sb (store byte), lui (load upper immediate). All memory access is through these explicit instructions.⁶
- **Control Flow:** beq (branch on equal), bne (branch on not equal), j (jump), jal (jump and link), jr (jump register). MIPS features a **branch delay slot**, meaning the instruction immediately following a branch or jump is always executed, regardless of whether the branch is taken.⁶ Compilers try to fill this slot with a useful instruction or a NOP.
- **Addressing Modes (MIPS is known for its simplicity here) ⁵:**
 - **Register:** Operand is in a GPR (e.g., add \$t1, \$t2, \$t3).
 - **Immediate:** Operand is a constant encoded in the instruction (e.g., addi \$t1, \$t2, 100).
 - **Displacement (Base + Offset):** Operand is in memory with address calculated as Base Register + 16-bit signed offset (e.g., lw \$t1, 32(\$s0)). This is the primary mode for memory access.⁶
 - **PC-relative:** Used for branch instructions, where the target address is PC + offset.
 - **Register Indirect (via displacement):** lw \$t0, 0(\$t1) is effectively register indirect addressing using register \$t1.

MIPS exemplifies a "pure" RISC approach with its fixed instruction length, limited addressing modes, and reliance on compilers for optimization. The load delay slot, while a performance feature in early non-interlocked pipelines, adds complexity for compilers and programmers.

5.3. ARM (Advanced RISC Machines)

- Overview and Philosophy:

ARM is a family of RISC ISAs that has become exceedingly dominant, particularly in mobile and embedded systems.⁹ Arm Holdings licenses the ISA and core designs. While rooted in RISC principles, ARM has evolved significantly, incorporating features to improve code density, performance, and power

efficiency, making it highly versatile.

- **Key Architectural Features (focus on ARMv7-A/AArch32 as a prominent example):**
 - **RISC Principles:** Load-store architecture, a large set of GPRs, and operations designed for efficient pipelining.⁹
 - **Word Size:** ARMv7-A is a 32-bit architecture (32-bit registers, 32-bit data paths). ARMv8-A introduced 64-bit capabilities (AArch64).⁹
 - **Registers (ARMv7-A A32 state):**
 - 16 directly accessible 32-bit GPRs (R0-R15). R13 is typically the Stack Pointer (SP), R14 is the Link Register (LR), and R15 is the Program Counter (PC).⁹
 - Current Program Status Register (CPSR) holds condition codes (N, Z, C, V), mode bits, interrupt disable flags.
 - **Instruction Sets & Length:**
 - **A32 (ARM state):** Fixed-length 32-bit instructions.⁹
 - **T32 (Thumb state):** Originally 16-bit instructions for better code density. Thumb-2 (introduced later) extended this to a mixed 16-bit and 32-bit variable-length instruction set, offering a balance of code density and performance.⁹
 - **A64 (ARMv8-A):** A new fixed-length 32-bit instruction set for 64-bit execution.⁹
 - **Conditional Execution:** A distinctive feature of A32 (and to some extent T32) is that most instructions can be executed conditionally based on the CPSR condition flags.⁹ A 4-bit condition field in the instruction determines this. This can reduce the need for short branches, improving performance and code density by avoiding pipeline flushes.¹⁵
 - **Architecture Profiles:** ARMv7 defines profiles like 'A' (Application, e.g., Cortex-A for smartphones), 'R' (Real-time, e.g., Cortex-R), and 'M' (Microcontroller, e.g., Cortex-M).⁹
- **Typical Instructions and Addressing Modes (A32):**
 - **Instruction Categories:** Data processing, data transfer (load/store), branch, coprocessor, exception-generating.
 - **Data Processing:** ADD, SUB, AND, ORR, EOR (XOR), MOV, CMP (compare), TST (test bits). Many can use a "barrel shifter" on one operand without an extra cycle, allowing shifts/rotates as part of the instruction (e.g., MOV R0, R1, LSL #2).
 - **Data Movement (Load/Store):** LDR (load register), STR (store register).

Supports various data sizes (byte, half-word, word). Multiple registers can be loaded/stored with LDM/STM.

- **Control Flow:** B (branch), BL (branch with link - saves return address in LR). Branches are typically conditional.
- **Addressing Modes (A32 offers more complex modes than classic MIPS):**
 - **Immediate Offset:** LDR R0, (Address = R1 + 4).
 - **Register Offset:** LDR R0, (Address = R1 + R2).
 - **Scaled Register Offset:** LDR R0, (Address = R1 + (R2 << 2)). This is powerful for array access.
 - **Pre-indexed:** LDR R0,! (Address = R1 + 4; R1 is updated to R1+4). The ! indicates write-back.
 - **Post-indexed:** LDR R0,, #4 (Address = R1; R1 is updated to R1+4 afterwards).
 - **PC-relative (Literals):** LDR R0, =label or LDR R0, [PC, #offset].
- **Floating-Point Instructions (A32/T32):** Most A32 floating-point instructions can be conditionally executed by appending a condition code suffix.¹⁶ However, certain floating-point instructions like VRINT, VSEL, VCVT (A,N,P,M forms), VMAXNM, VMINNM cannot be used in an IT (If-Then) block in T32. Specifying other floating-point instructions in an IT block is also deprecated.¹⁶ The IT instruction provides conditional execution for a small block of subsequent T32 instructions.

ARM's evolution demonstrates a pragmatic approach to RISC design. While maintaining core RISC tenets, it has incorporated features like conditional execution and the Thumb instruction set to address practical needs like code density and performance in diverse applications, contributing significantly to its widespread adoption. The shift in AArch64 to reduce general conditional execution in favor of specific conditional select instructions (like CSEL) reflects evolving trade-offs in processor design, where more complex branch predictors and the overhead of encoding conditions on every instruction became factors.¹⁵

5.4. Comparative Summary

The following table provides a high-level comparison of these three ISAs:

| Feature | LC-3 | MIPS (MIPS32 I) | ARM (ARMv7-A) |
|---------|------|-----------------|---------------|
|---------|------|-----------------|---------------|

| | | | |
|---------------------------------|---|--|--|
| | | | A32) |
| Primary Design Goal | Educational, simplicity ¹³ | Classic RISC, performance, academic ⁶ | Versatile RISC, mobile/embedded, power efficiency ⁹ |
| RISC/CISC Classification | RISC-like ¹² | RISC ⁶ | RISC (with pragmatic extensions) ⁹ |
| Word Size | 16-bit ¹³ | 32-bit ⁶ | 32-bit ⁹ |
| GPRs (Number & Size) | 8 x 16-bit ¹³ | 32 x 32-bit (plus HI/LO) ⁶ | 16 x 32-bit (some specialized) ⁹ |
| Instruction Length(s) | 16-bit fixed ¹³ | 32-bit fixed ⁶ | 32-bit fixed (A32), 16/32-bit (T32) ⁹ |
| Key Addressing Modes | PC-relative, Base+offset, Indirect, Immediate ¹² | Register, Immediate, Base+Displacement ⁵ | Register, Immediate, Pre/Post-indexed (offset/scaled register), PC-relative ⁹ |
| Notable Features | Simplicity, Condition Codes (N,Z,P) ¹² | Branch Delay Slot, \$0 register ⁶ | Conditional Execution (most instructions), Barrel Shifter, Thumb ISA ⁹ |
| Typical Applications | Teaching computer architecture ¹³ | Embedded systems, networking, education ⁵ | Mobile phones, tablets, embedded systems, servers (ARMv8+) ⁹ |

This comparison highlights how different ISAs are tailored for different purposes, from the pedagogical clarity of LC-3 to the streamlined RISC principles of MIPS, and the highly successful and adaptable RISC design of ARM. Each makes different trade-offs

in terms of complexity, instruction set richness, and features to meet its specific goals.

6. Concluding Remarks and Synthesis

The Instruction Set Architecture is undeniably the bedrock upon which software and hardware interaction is built. As the definitive contract between these two realms, the ISA dictates the processor's fundamental capabilities and the language through which software commands the hardware.¹ The design of an ISA involves a complex tapestry of trade-offs: the richness of the instruction set versus the simplicity of hardware implementation (the CISC vs. RISC dilemma⁷); the efficiency of fixed-length instructions versus the code density of variable-length instructions⁴; the number of explicit operands per instruction impacting instruction length and register pressure³; and the variety of addressing modes influencing programming flexibility and hardware complexity.¹⁰

These choices have profound implications for system performance, power consumption, cost, and the entire software development ecosystem. An ISA optimized for a particular domain—such as ARM's low-power designs for mobile and embedded systems⁹, or specialized extensions for digital signal processing or cryptography—can yield significant advantages over general-purpose ISAs within that niche. The longevity and success of an ISA often depend on its ability to foster a robust ecosystem of tools, compilers, operating systems, and skilled developers, as seen with architectures like x86 and ARM.

The landscape of ISA development is continually evolving, driven by new computational demands and technological advancements. Several key trends are shaping its future:

- **Domain-Specific Architectures (DSAs) and Extensions:** There is a growing movement towards ISAs or ISA extensions tailored for specific workloads, such as Artificial Intelligence/Machine Learning (e.g., Google's TPUs, NVIDIA's tensor cores), graphics processing (GPUs), and network processing. The ability to add custom instructions, as Arm allows for its Cortex-M cores to accelerate specialized workloads², is indicative of this trend. This specialization can offer orders-of-magnitude improvements in performance and energy efficiency for targeted tasks compared to general-purpose CPUs.
- **Open ISAs:** The rise of open-source ISAs, most notably RISC-V, is challenging traditional proprietary ISA models. Open ISAs promote collaboration, innovation, and customization, allowing designers to create tailored processors without

licensing fees and with greater transparency. This can lower barriers to entry for hardware development and foster a diverse ecosystem of specialized cores.

- **Enhanced Security Features:** With increasing cybersecurity threats, ISAs are incorporating more robust security features directly into the hardware. This includes mechanisms for memory protection, secure enclaves (e.g., ARM TrustZone), hardware-accelerated cryptography, and features to mitigate side-channel attacks.
- **Continued Focus on Energy Efficiency:** As mobile devices, IoT, and edge computing proliferate, the demand for ultra-low-power processing remains paramount. ISA designs will continue to prioritize energy efficiency through simpler instructions, power-aware microarchitectures, and features that allow fine-grained power management.
- **Modular ISAs:** Modern ISAs are increasingly designed with a standard base and a set of optional extensions (e.g., MIPS extensions⁶, RISC-V's modular approach). This allows implementers to select and include only the features relevant to their specific application, optimizing for cost, power, and area, while maintaining compatibility with a common software base.

A persistent theme in ISA evolution is the balance between standardization and specialization. Standardization, as exemplified by the broad adoption of x86 and ARM, cultivates vast software ecosystems and ensures widespread compatibility.² However, the relentless pursuit of performance and efficiency for emerging workloads like AI is driving the need for specialized hardware. Modular ISAs represent an attempt to reconcile these forces, offering a common foundation that can be extended for specific needs. The future likely involves a heterogeneous computing landscape where general-purpose cores coexist with an increasing number of specialized accelerators, all underpinned by carefully designed ISAs that continue to bridge the ever-evolving worlds of hardware and software.

Works cited

1. [www.arm.com](https://www.arm.com/glossary/isa#:~:text=An%20Instruction%20Set%20Architecture%20(ISA.as%20how%20it%20gets%20done.), accessed June 4, 2025,
[https://www.arm.com/glossary/isa#:~:text=An%20Instruction%20Set%20Architecture%20\(ISA.as%20how%20it%20gets%20done.](https://www.arm.com/glossary/isa#:~:text=An%20Instruction%20Set%20Architecture%20(ISA.as%20how%20it%20gets%20done.)
2. What is Instruction Set Architecture (ISA)? – Arm®, accessed June 4, 2025,
<https://www.arm.com/glossary/isa>
3. Instruction Formats in Computer Organization: Types, Examples, accessed June 4, 2025,
<https://www.ccbp.in/blog/articles/instruction-formats-in-computer-organization>

4. Machine Level Instructions (in the General Model) - Teaching, accessed June 4, 2025,
https://teaching.idallen.com/dat2343/11w/notes/410_MachineLevelInstructions.html
5. Instruction Set Architecture - CS2100 - NUS Computing, accessed June 4, 2025,
<https://www.comp.nus.edu.sg/~adi-yoga/CS2100/ch07/>
6. MIPS architecture - Wikipedia, accessed June 4, 2025,
https://en.wikipedia.org/wiki/MIPS_architecture
7. What Is RISC and CISC Architecture & Their Differences | Glossary ..., accessed June 4, 2025,
<https://conclusive.tech/glossary/what-is-risc-and-cisc-architecture-their-differences/>
8. RISC vs CISC: Key Differences - DiffStudy, accessed June 4, 2025,
<https://diffstudy.com/risc-vs-cisc/>
9. ARM architecture family - Wikipedia, accessed June 4, 2025,
https://en.wikipedia.org/wiki/ARM_architecture_family
10. Addressing Modes | GeeksforGeeks, accessed June 4, 2025,
<https://www.geeksforgeeks.org/addressing-modes/>
11. A Quick Guide to Addressing Modes in Computer Architecture - Board Infinity, accessed June 4, 2025,
<https://www.boardinfinity.com/blog/a-quick-guide-to-addressing-modes/>
12. Simplified LC-3 Instruction set - Coert Vonk, accessed June 4, 2025,
<https://coertvonk.com/inquiries/how-cpu-work/instruction-set-30971>
13. Little Computer 3 - Wikipedia, accessed June 4, 2025,
https://en.wikipedia.org/wiki/Little_Computer_3
14. ARM Architecture Reference Manual ARMv7-A and ARMv7-R edition, accessed June 4, 2025,
<https://developer.arm.com/documentation/ddi0406/c/Application-Level-Architecture/Instruction-Details/Conditional-execution>
15. Why are conditionally executed instructions not present in later ARM instruction sets?, accessed June 4, 2025,
<https://stackoverflow.com/questions/22168992/why-are-conditionally-executed-instructions-not-present-in-later-arm-instruction>
16. Arm Instruction Set Reference Guide - Arm Developer, accessed June 4, 2025,
<https://developer.arm.com/documentation/100076/0100/Advanced-SIMD-and-Floating-point-Programming/Floating-point-Programming/Conditional-execution-of-A32-T32-floating-point-instructions?lang=en>